



Contemporary Architecture

Chapter 7



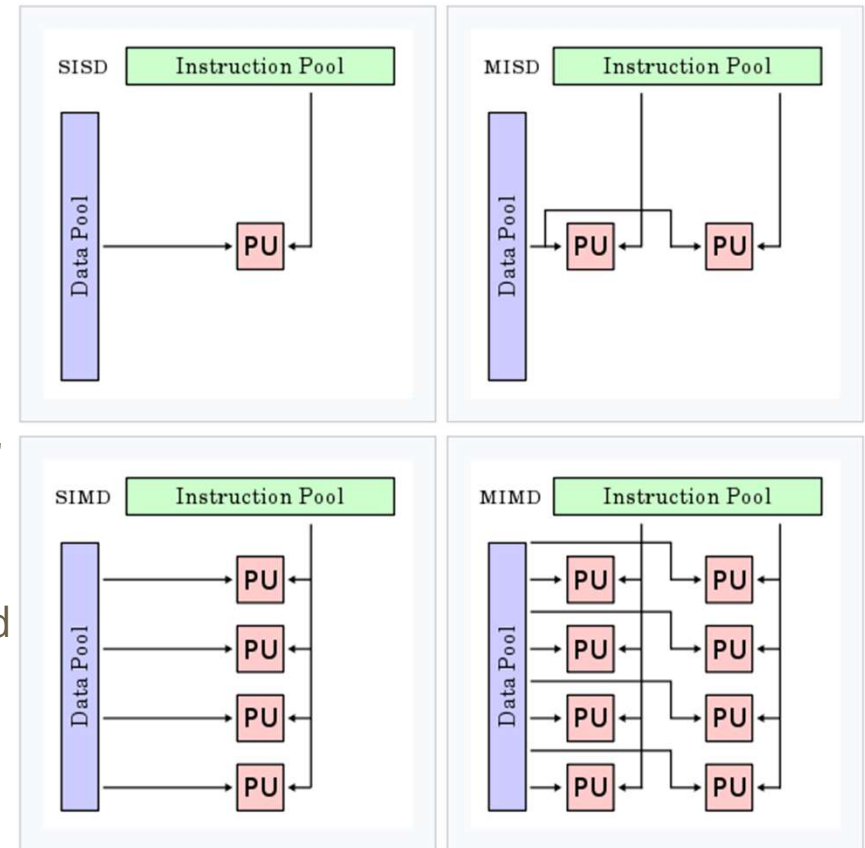
Contemporary Architecture

- ★ Flynn's Taxonomy
 - SISD, SIMD, MISD, MIMD
- ★ Instruction-Level Parallelism
 - Vector Processor, Vector Unit
 - Pipelining Processor, SuperPipeline
 - SuperScalar (Out-of-order execution, Register Renaming)
 - Very Long Instruction Word (VLIW)
- ★ SuperSymmetry
- ★ Dynamic vs. Static Scheduling
- ★ Compiler & ILP
 - Loop Unroll, Software Pipelining, LLVM
- ★ Thread-Level Parallelism



Flynn's Taxonomy

- ★ by Michael Flynn (1934-) in 1966
- ★ Partition by Instruction Stream and Data Stream
 - Single Instruction stream, Single Data stream (SISD) - **uniprocessor**
 - Single Instruction stream, Multiple Data stream (SIMD) - **vector processor**
 - Multiple Instruction stream, Single Data stream (MISD) - **Fault-Tolerant Computing** (The course is available as approved from A. Arthit.)
 - Multiple Instruction stream, Multiple Data stream (MIMD) - **Multicore, Distributed Systems**



Pictures from wikipedia



Instruction-Level Parallelism

- ★ ILP - a measure of how many of the instructions in a computer program can be executed simultaneously
- ★ Several micro-architectural techniques involve:
 - Vectorization
 - Speculative Execution - allow (parts of) instructions to be executed before needed.
 - Branch Prediction - avoid stalling from control dependencies(speculative execution)
 - Instruction Pipelining and Super Pipeline
 - SuperScalar
 - Out-of-order execution
 - Register Renaming
 - VLIW



Vector Processor, Vector Unit

- ★ A variation of SIMD
- ★ Known products
 - Apple/Motorola/IBM AltiVec
 - Intel MMX, SSE[2,3,4], AVX[2,512]
 - GPU - CUDA, OpenCL
- ★ Apply single operation to vectors (array of data)
- ★ We will revisit this on GPGPU topics.

```
// uniprocessor
int A[8];
int B[8];
int C[8];

for (int i=0;i<8;i++)
{
    C[i]=A[i]+B[i];
}
```

```
// vectorize code
// equivalent in C
int A[8];
int B[8];
int C[8];

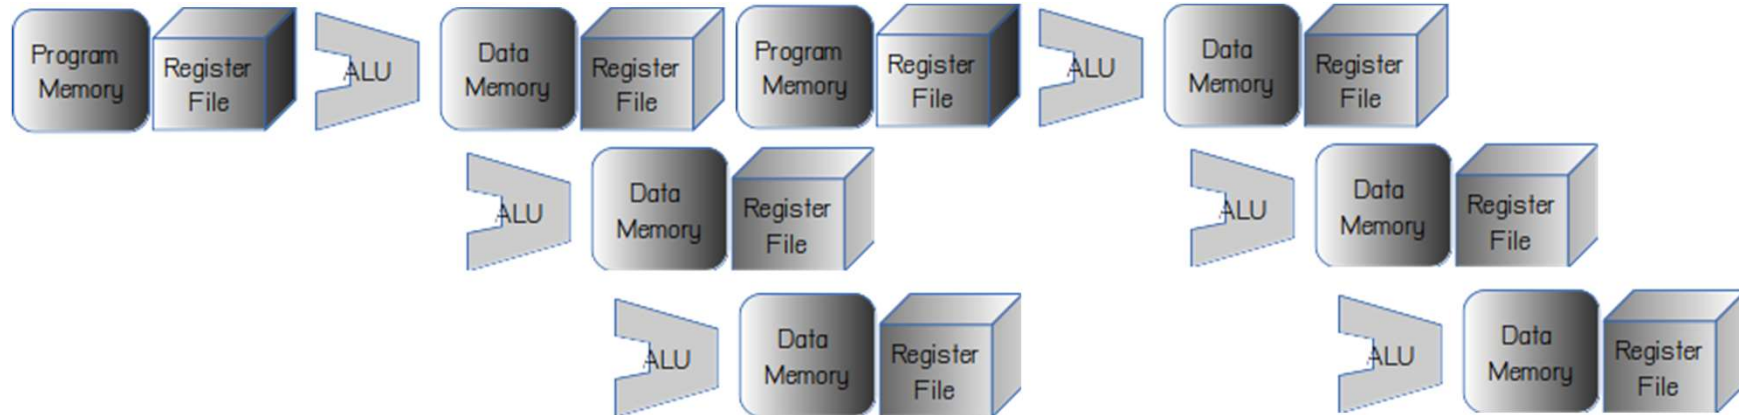
C=A+B;
```

```
; Optimize of -O3 on GCC 6.3
; vectorize with Intel AVX
...
movdqa    16+B(%rip), %xmm0
paddb     16+A(%rip), %xmm0
movaps    %xmm0, 16+C(%rip)
...
```



Vector Processor (ctd.)

- ★ Some vector processors will expand loops internally.





Vector processor in real world

- ★ Ordering several dishes of the same things, a cook receives an order once.



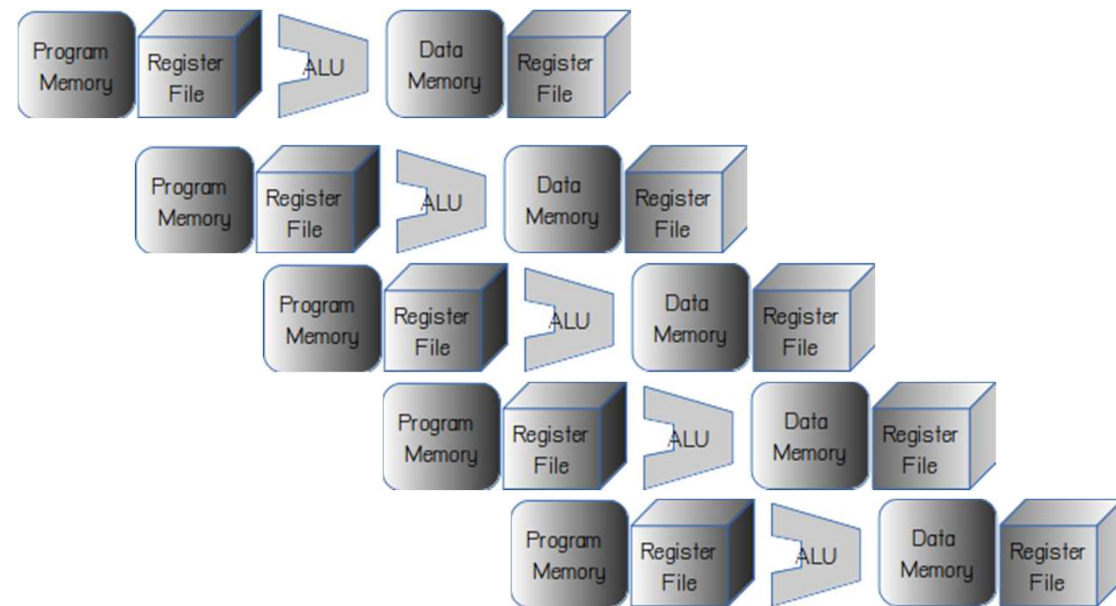
I want 10 cups of noodle.





Pipelining Processor (revisit)

- ★ A pipeline with 5 states has 5 instructions in flight.
- ★ For k state pipelines, $ILP = k$.
- ★ CPU Time for n instructions
 $T(n) = (k + (n-1)) T_c$





Pipeline in real world

- ★ One line for buying food.
- ★ 5 steps





SuperScalar


- ★ SuperScalar with degree 2 and 5 states of pipeline has 10 instructions in flight.
- ★ For k state pipelines in superscalar x degrees, $ILP = kx$.
- ★ CPU Time for n instructions $T(n) = (k + (n-x)/x) T_c$





SuperScalar in real world

- ★ One line for buying food.
- ★ After getting in, a manager will separate customers into two lines (rice line and noodle line)





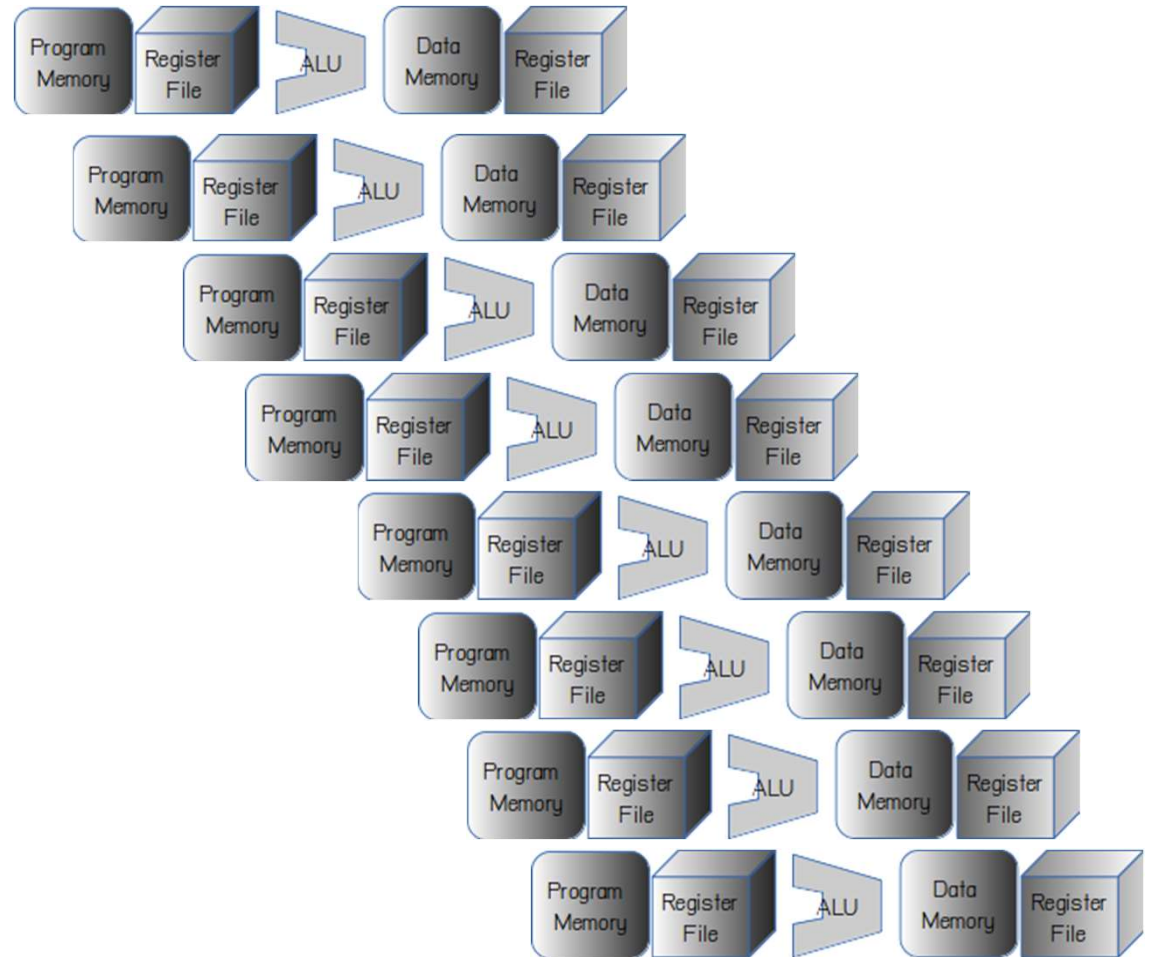
SuperScalar (ctd.)

- ★ With super scalar, we may have **out-of-order execution**.
- ★ A person that comes after (but order noodle) may finish earlier than a person that comes before (but order rice).
- ★ This may result in potential **WAR** and **WAW**.
- ★ To solve this -- **in-order completion** ?



SuperPipeline


- ★ Use minor cycles.
- ★ In this example, 2 minor cycles is equivalent to 1 major cycle.
- ★ For m minor cycles to k state pipeline,
 $ILP = mk$
- ★ CPU Time for n instructions
 $T(n) = (k + (n-1)/m) T_c$





SuperPipeline SuperScalar

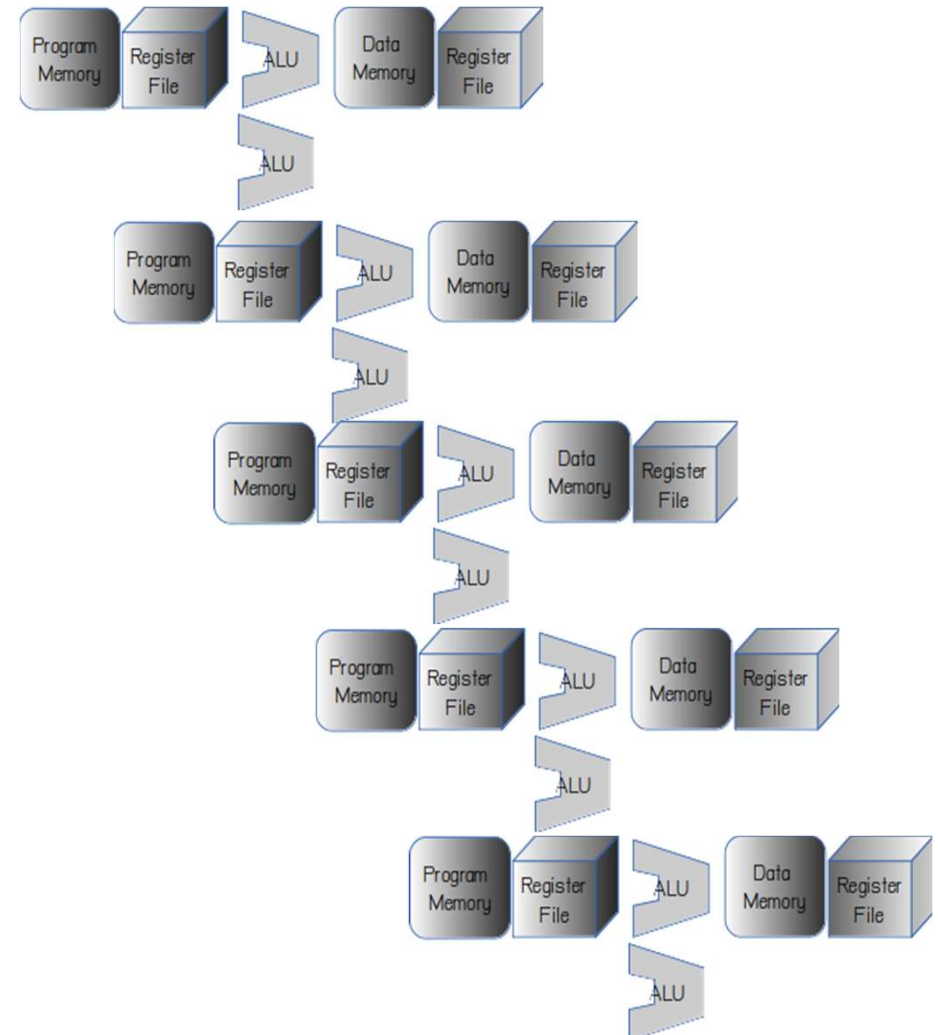
- ★ SuperScalar with minor cycles.
- ★ In this example, 2 minor cycles is equivalent to 1 major cycle.
- ★ For m minor cycles to k states pipeline in superscale degree x,
ILP = mkx
- ★ CPU Time for n instructions
 $T(n) = (k + (n-x)/mx) T_c$






VLIW

- ★ VLIW with degree 2 and 5 states of pipeline has 10 instructions in flight.
- ★ For k state pipelines in VLIW x degrees,
ILP = kx.
- ★ CPU Time for n major instructions
 $T(n) = (k + (n/x - x)) T_c$





VLIW in real world



- ★ Static scheduling.
- ★ Two lines for ordering
- ★ One for rice
- ★ Another for noodle



SuperSymmetry

★ SuperScalar

★ Speedup=
$$\frac{(k + (n-x)/x) T_c}{(k + (n-1)) T_c}$$

★ For large n,
Speedup=x

★ SuperPipeline

★ Speedup=
$$\frac{(k + (n-1)/m) T_c}{(k + (n-1)) T_c}$$

★ For large n,
Speedup=m

★ VLIW

★ Speedup=
$$\frac{(k + (n/x-x)) T_c}{(k + (n-1)) T_c}$$

★ For large n,
Speedup=x

At the same degree, SuperScalar, SuperPipeline and VLIW yield the same performance.



Observations on VLIW

- ★ Software for Intel Itanium and Intel Itanium II (also the successors) are not binary compatible. **Why?** (Hint. static scheduling)
- ★ As of January 2019, Intel officially announced the end of life and product discontinuance of the Itanium CPU family.
- ★ Other implementations of VLIW exist (e.g. Transmeta Crusoe). They are not practically used today.





Speed Up (pipeline)

- ★ Assuming that there are two separate pipelines (integer unit and floating point unit), two instructions can be executed at the sametime.

- ★ CPU Time = Max (Int, FP)

- ★ Potential Speed Up

$$= 19 / \text{Max}(\text{Int}, \text{FP})$$

$$= 19 / \text{Max}(7, 12)$$

$$= 1.6$$

Basic Loop:

	Cycles
load Ra <- Ai	1
load Ry <- Yi	1
fmult Rm <- Ra * Rx	7
fadd Rs <- Rm + Ry	5
store Ai <- Rs	1
inc Yi	1
dec i	1
inc Ai	1
branch	1



Performance (Revisit)

- ★ CPU Time = IC x CPI x TC
- ★ In practical, operation classes and latency are not balance.

	IC	CPI	TC	Potential Speed Up
Single Cycle Processor	n	1	L	1
Multiple Cycle Processor	n	3-5 (avg=4.2)	L/5	1.xx
Pipelining Processor	n	1	L/5	5
SuperScalar (degree k)	n	1/k	L/5	5k
SuperPipeline (m minor cycle)	n	1/m	L/5	5m
VLIW (x instructions in a word)	n/x	1	L/5	5x



Compiler & ILP

- ★ Compiler can play a critical role in exploiting more ILP
- ★ Small basic block means less opportunity for filling available slots.
- ★ A larger basic block means more opportunities for filling slots.
- ★ To increase a basic block,
 - Loop Unrolling
 - Software Pipeline



Loop Unroll

- ★ Make the basic block in a loop larger by having more iterations in a loop.
- ★ The new code is potentially not stall at the cost of larger code size.

Normal loop.

```
for (i=MAX-1;i>=0;i--) {  
    A[i] = A[i] + C;  
}
```

Symbolically unroll 4 iterations

```
for (i=MAX-1;i>=3;i-=4) {  
    A[i] = A[i] + C;  
    A[i-1] = A[i-1] + C;  
    A[i-2] = A[i-2] + C;  
    A[i-3] = A[i-3] + C;  
}
```



Loop Unroll (ctd.)

Normal Loop

```

LOOP:
    LW      $f0, $r1
(stall)
    ADD     $f2, $f0, $f1
(stall)
(stall)
    SW      $f2, $r1
    SUB     $r1, $r1, #4
    BNE     $r1, LOOP
  
```

8 cycles / Iteration

Register Renaming

- ★ To free unrolled iterations from the first iteration, a register would be renamed.
- ★ Each iteration is now resource independent.

Unroll (4 iterations)

```

LOOP:
    LW      $f0, 0($r1)
(stall)
    ADD     $f2, $f0, $f1
(stall)
(stall)
    SW      $f2, 0($r1)
    LW      $f3, 4($r1)
(stall)
    ADD     $f4, $f3, $f1
(stall)
(stall)
    SW      $f4, 4($r1)
    LW      $f5, 8($r1)
(stall)
    ADD     $f6, $f5, $f1
(stall)
(stall)
    SW      $f6, 8($r1)
    LW      $f7, 12($r1)
(stall)
    ADD     $f8, $f7, $f1
(stall)
(stall)
    SW      $f8, 12($r1)
    SUB     $r1, $r1, #16
    BNE     $r1, LOOP
  
```

Unroll + Rescheduling

```

LOOP:
    LW      $f0, 0($r1)
    LW      $f3, 4($r1)
    LW      $f5, 8($r1)
    LW      $f7, 12($r1)
    ADD     $f2, $f0, $f1
    ADD     $f4, $f3, $f1
    ADD     $f6, $f5, $f1
    ADD     $f8, $f7, $f1
    SW      $f2, 0($r1)
    SW      $f4, 4($r1)
    SW      $f6, 8($r1)
    SW      $f8, 12($r1)
    SUB     $r1, $r1, #16
    BNE     $r1, LOOP
  
```

14 cycles / 4 Iteration
= 3.5 cycles per iteration

Speed Up=8/3.5



Loop Unroll and Superscalar

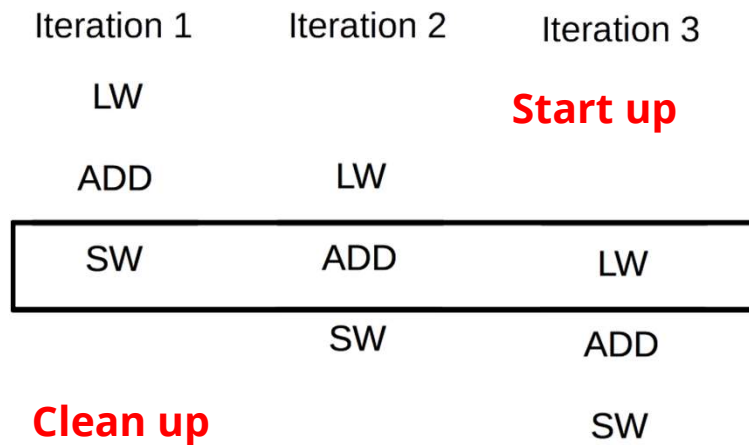
12 cycles / 5 Iteration
= 2.4 cycles per iteration

Integer Unit	Floating Unit	Clock Cycle (Issued)
LOOP:		
LW \$f0 , 0(\$r1)		1
LW \$f3, 4(\$r1)		2
LW \$f5, 8(\$r1)	ADD <u>\$f2</u> , \$f0 , \$f1	3
LW \$f7, 12(\$r1)	ADD \$f4, \$f3, \$f1	4
LW \$f9, 16(\$r1)	ADD \$f6, \$f5, \$f1	5
SW <u>\$f2</u> , 0(\$r1)	ADD \$f8, \$f6, \$f1	6
SW \$f4, 4(\$r1)	ADD \$f10, \$9, \$f1	7
SW \$f6, 8(\$r1)		8
SW \$f8, 12(\$r1)		9
SW \$f10, 16(\$r1)		10
SUB \$r1, \$r1, #24		11
BNE \$r1, LOOP		12



Software Pipeline

- ★ Unroll as needed only.
- ★ Good performance, appropriate code size.



Unroll (3 iterations)

```

LOOP:
    LW    $f0, 0($r1)
    LW    $f3, 4($r1)
    LW    $f5, 8($r1)
    ADD   $f2, $f0, $f1
    ADD   $f4, $f3, $f1
    ADD   $f6, $f5, $f1
    SW    $f2, 0($r1)
    SW    $f4, 4($r1)
    SW    $f6, 8($r1)
    SUB   $r1, $r1, #12
    BNE   $r1, LOOP
    
```

Unroll + Rescheduling

Start up

Clean up

```

SUB    $r1, $r1, #4
LW     $f0, 4($r1)
ADD    $f2, $f0, $f1
LW     $f0, 0($r1)
LOOP:
    SW    $f2, 0($r1)
    ADD   $f2, $f0, $f1
    LW    $f0, 8($r1)
    SUB   $r1, $r1, #4
    BNE   $r1, LOOP

    SW    $f2, 4($r1)
    ADD   $f2, $f0, $f1
    SW    $f2, 0($r1)
    
```



Dynamic vs. Static Scheduling

★ **Dynamic Scheduling (Hardware)**

- ★ Hardware dynamically schedules code (out-of-order execution).
 - Loop are dynamically unrolled in the processor.
 - Multiple instructions are dynamically issued.
- ★ Legacy software gains better performance on new hardware.
- ★ Complex hardware design?

★ **Static Scheduling (Software/Compiler)**

- ★ Compiler can see larger scope.
 - Loop are statically unrolled by compilers.
 - Multiple instructions are statically scheduled for each cycle (see VLIW).
- ★ Poor compatibilities
- ★ Dependency cannot be determined at compile time.
- ★ Simple hardware
- ★ Lower Power requirements?

Who is the winner?



How about a chip with multicores?





Chip with multicores

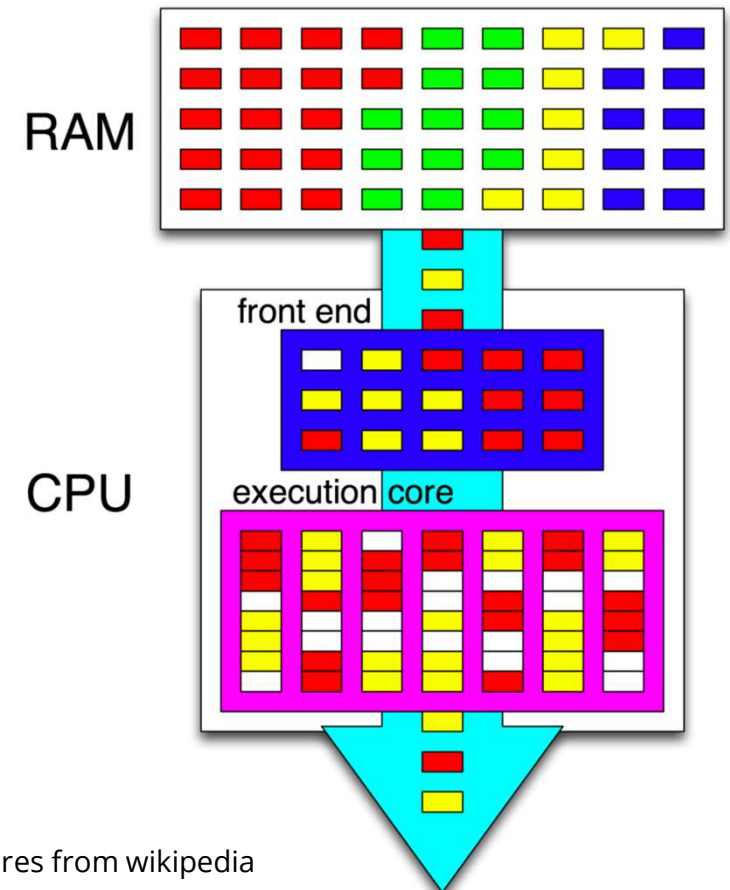
- ★ A chip with multicores is basically a chip with multiple superscalar cores inside.
- ★ Can also be viewed as Multiple Pipes.
- ★ Demands for more ILP. -- Thread-Level parallelism



Thread-Level Parallelism

- ★ Since a program (thread) cannot fill available slots in the processor, let's get instructions from several threads.
- ★ Instructions in pipeline are coming from several threads.
- ★ The technology is called simultaneous multithreading (SMT).
- ★ Intel implementation is officially called Hyper-Threading Technology.

The detail is beyond the scope of this course. To learn more, join **High Performance Architecture**.



Pictures from wikipedia



Summary

- ★ There exist several approaches to ILP.
- ★ At similar degrees, SuperScalar, SuperPipeline, VLIW will yield similar performance.
- ★ Dynamic scheduling allows legacy software to gain performance.
- ★ Static scheduling may gain more benefits if a compiler is more powerful.
 - Proven to be too difficult to have good compiler.



Exercises





Loop Unroll

- ★ For the given code, if MAX is a known/large value,
 - (1) how many iterations can we unroll this loop?
 - (2) What is the potential speed up of the unrolled (rescheduled) code comparing to the original one.
 - (3) What will be the size of the unrolled code?
 - (4) Which resources is limited to the number of iterations that can be unrolled?

```
for (i=MAX-1; i>=0; i--) {  
    A[i] = A[i] + C;  
}
```

```
LOOP:  
    LW    $f0,  
    $r1  
    (stall)  
    ADD   $f2, $f0, $f1  
    (stall)  
    (stall)  
    SW    $f2, $r1  
    SUB   $r1, $r1, #4  
    BNE   $r1, LOOP
```




Software Pipeline

- ★ For the given code, if MAX is a known/large value, please use software pipeline to optimize the code.
 - (1) What is the potential speed up of the unrolled (rescheduled) code comparing to the loop unroll version.
 - (2) What will be the size of the code?

```
for (i=MAX-1; i>=0; i--) {  
    A[i] = A[i] + C;  
}
```

```
LOOP:  
    LW    $f0,  
    $r1  
    (stall)  
    ADD   $f2, $f0, $f1  
    (stall)  
    (stall)  
    SW    $f2, $r1  
    SUB   $r1, $r1, #4  
    BNE   $r1, LOOP
```



Dynamic SuperScalar (out-of-order execution)

- ★ Given a super scalar at degree two, please draw a diagram showing the pipeline of the following code.

```
LOOP:
    LW    $f0, $r1
    ADD   $f2, $f0,$f1
    SW    $f2, $r1
    SUB   $r1, $r1,#4
    BNE   $r1, LOOP
```

		Issue	EX	MEM	WB	Notes
LW	\$f0, \$r1	1	2	3	4	
ADD	\$f2, \$f0,\$f1	1	5,6,7		8	Wait \$f0
SW	\$f2, \$r1	2	3	9		Wait \$f2
SUB	\$r1, \$r1,#4	2	4		5	
BNE	\$r1, LOOP	3	6			
LW	\$f0, \$r1	4	7	8	9	
ADD	\$f2, \$f0,\$f1	4	10,11,12		13	Wait \$f0
SW	\$f2, \$r1	5	8	14		
SUB	\$r1, \$r1,#4	5	9		10	
BNE	\$r1, LOOP	6	11			



End of Chapter 7

