



Memory Organization (part 2)

Chapter 8



Memory Organization (part 2)

- ★ Virtual Memory
- ★ Paging Organization
- ★ Address Translation
- ★ Page Tables
 - Translation Look-Aside Buffers
- ★ Virtual Memory and Cache
 - Overlapped Cache & TLB Access



Virtual Memory

- ★ Cost (Run software that requires 4GB on 2GB hardware.)
 - Illusion of more memory (borrowing from secondary storage)
 - Demand loading and swap
- ★ Program Relocatable
 - Different process has its own address space.
 - Address 0x100 of Prog. A is not the same as address 0x100 of Prog. B.
- ★ Protection (Security)
 - Prog. A may not access memory space of Prog. B directly.



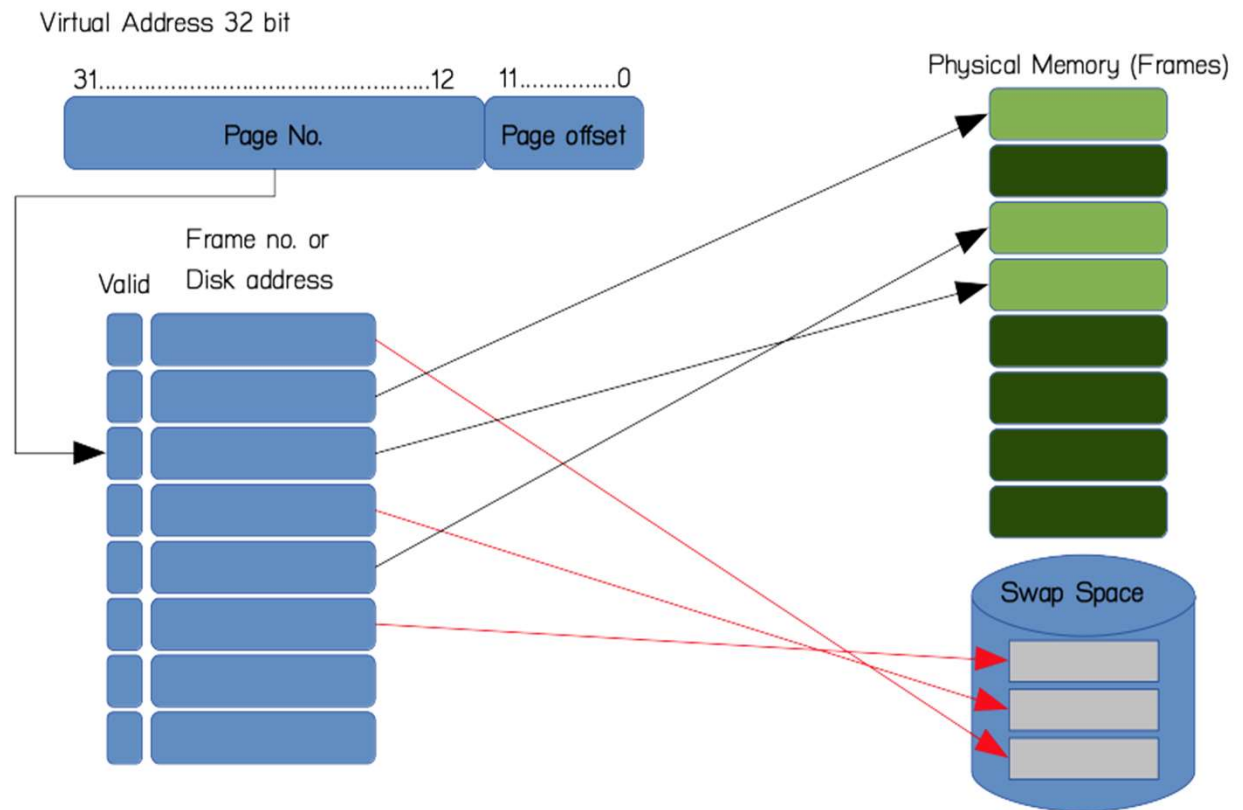
Paging Organization

- ★ A virtual address is the address that visible to software.
- ★ A virtual address is translated into a physical address by looking up from page table.
- ★ When physical address space is smaller than the amount of virtual address space, a swap space is borrowed from secondary storage (disk).
- ★ A virtual block is called a page. (Similar to block in cache)
- ★ A physical block is called a page frame.
- ★ A miss is called page faults.



Paging Organization (ctd.)

- ★ Physical Memory can act as a cache for secondary storage (file/disk).
- ★ Missing item is loaded on fault only (demand load policy)
- ★ Replacement policy is handled by operating system.





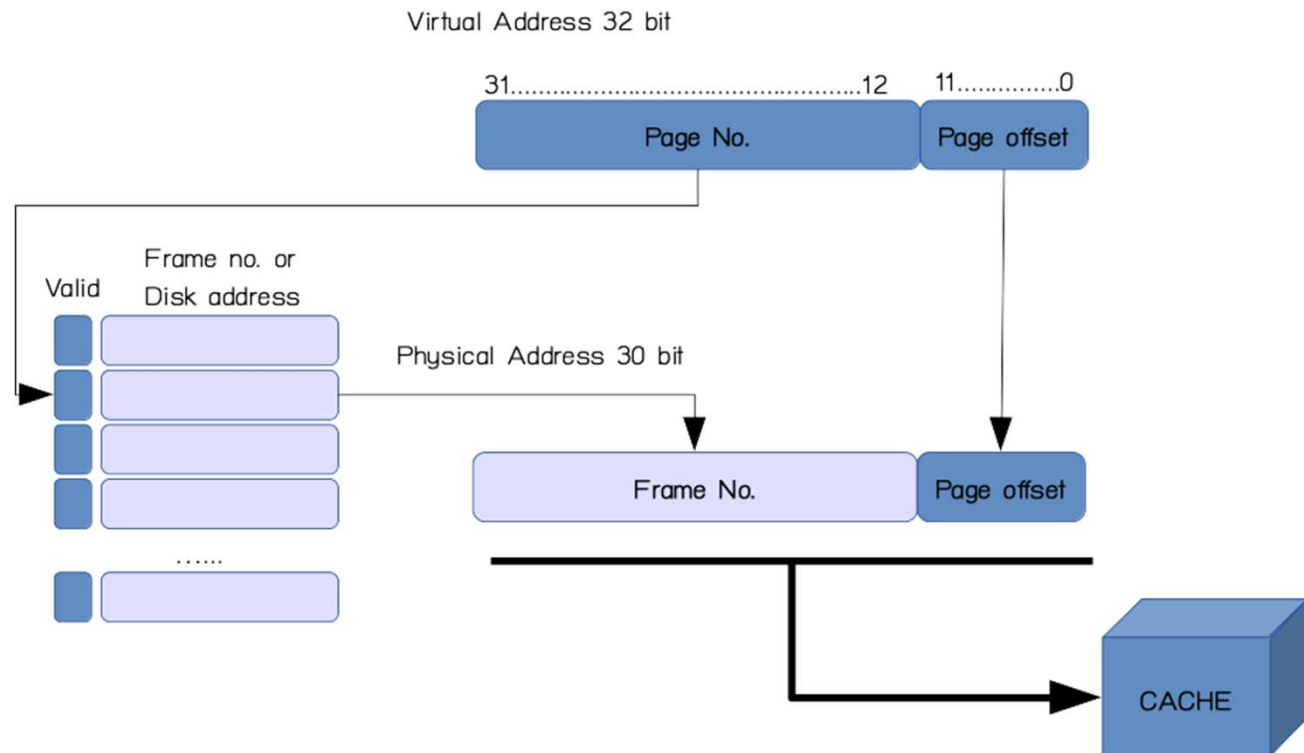
Pages: Virtual Memory Blocks

- ★ Virtual and physical address space partitioned into blocks of equal size
- ★ Virtual -- pages, Physical -- page frames
- ★ Page faults -- data is not in memory, retrieve from disk
 - High miss penalty. A page is usually large (eg. 4KB).
 - A write through is too expensive. Only uses **write back** (page out).
 - A fault is handled by software.
(Hardware does not know file system or swap management.)



Address Translation

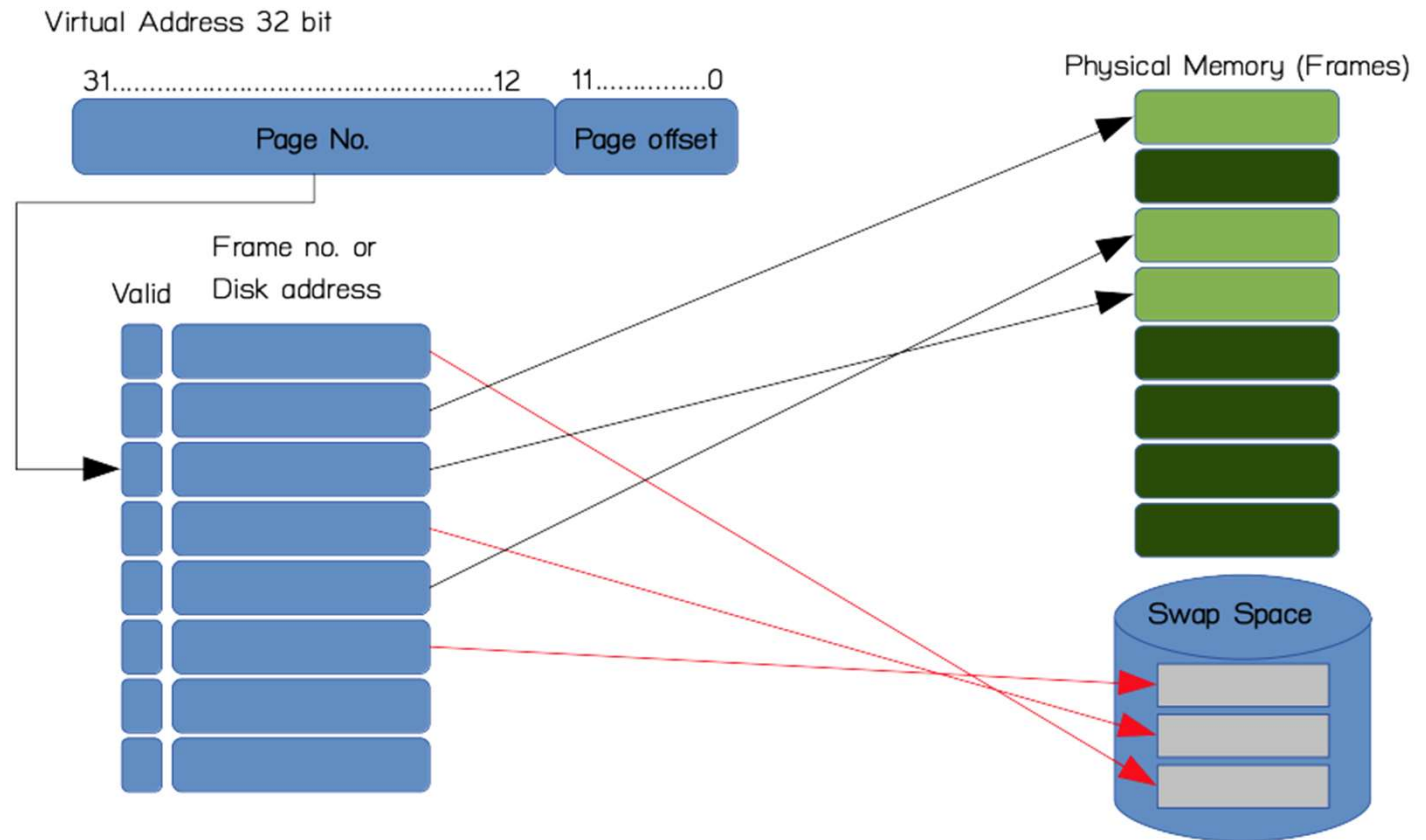
- ★ Virtual address space can be larger than Physical address space. (Run 4GB software on 2GB hardware)





Page Tables


- ★ In reality, page table also include access permission (eg. readonly)





Page Table (ctd.)

- ★ From this organization, what is the size of :
 - Page
 - Physical memory
 - A page table (4 bytes x 2^{20})
- ★ A page table is per process

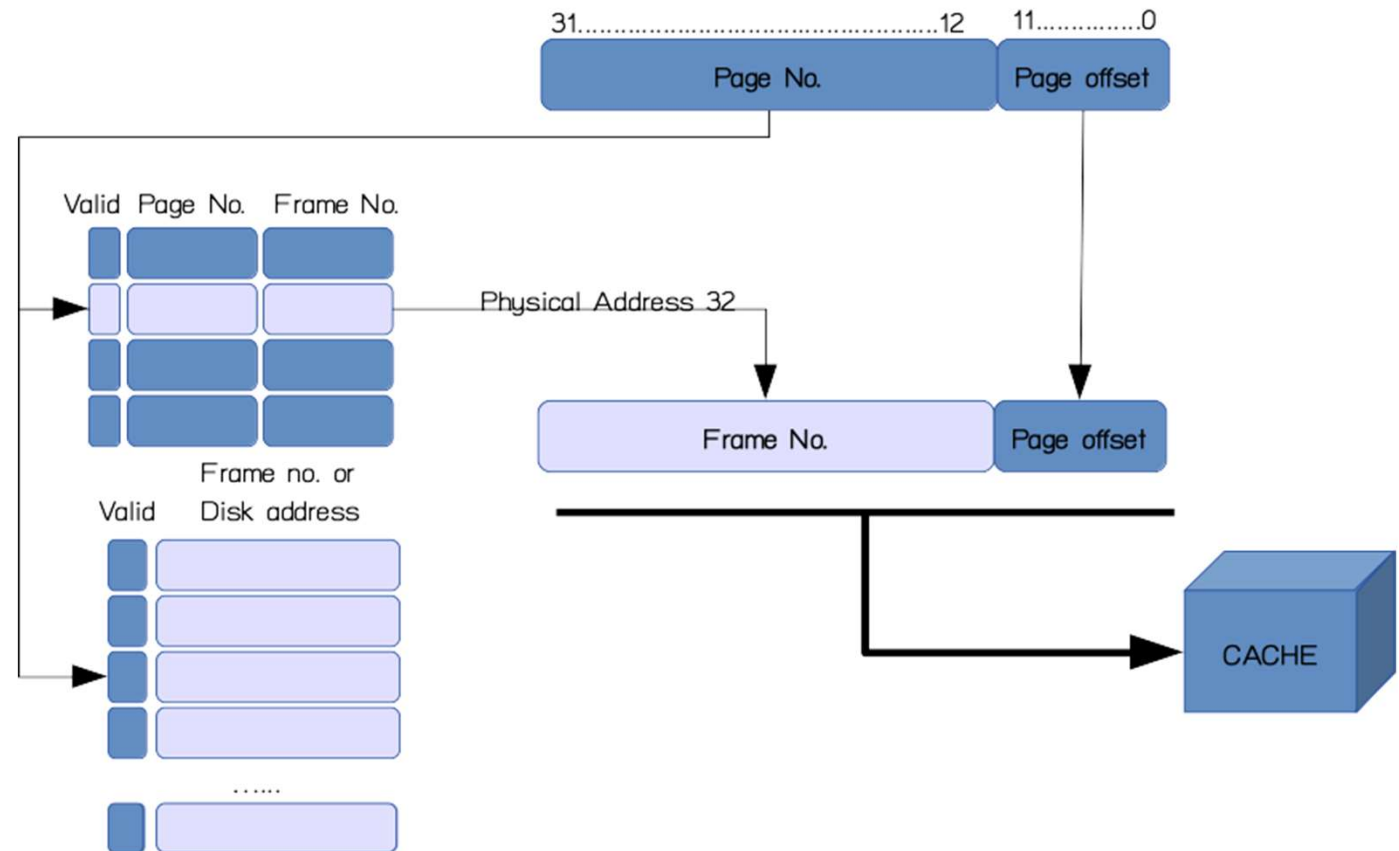




Make Address Translation fast

Virtual Address 32 bit

- ★ A cache for address translations: translation lookaside buffer





TLB (ctd.)


- ★ A way to speed up translation is to use a special cache of recently used page table entries -- this has many names, but the most frequently used is Translation Lookaside Buffer or TLB
- ★ TLB access time comparable to cache access time (much less than main memory access time)

Virtual Address	Physical Address	Dirty	Ref	Valid	Access



Translation Look-Aside Buffers

- ★ TLBs are usually small, typically not more than 128 - 256 entries even on high-end machines. This permits **fully associative** lookup on these machines. Most mid-range machines use **small n-way set associative** organizations.

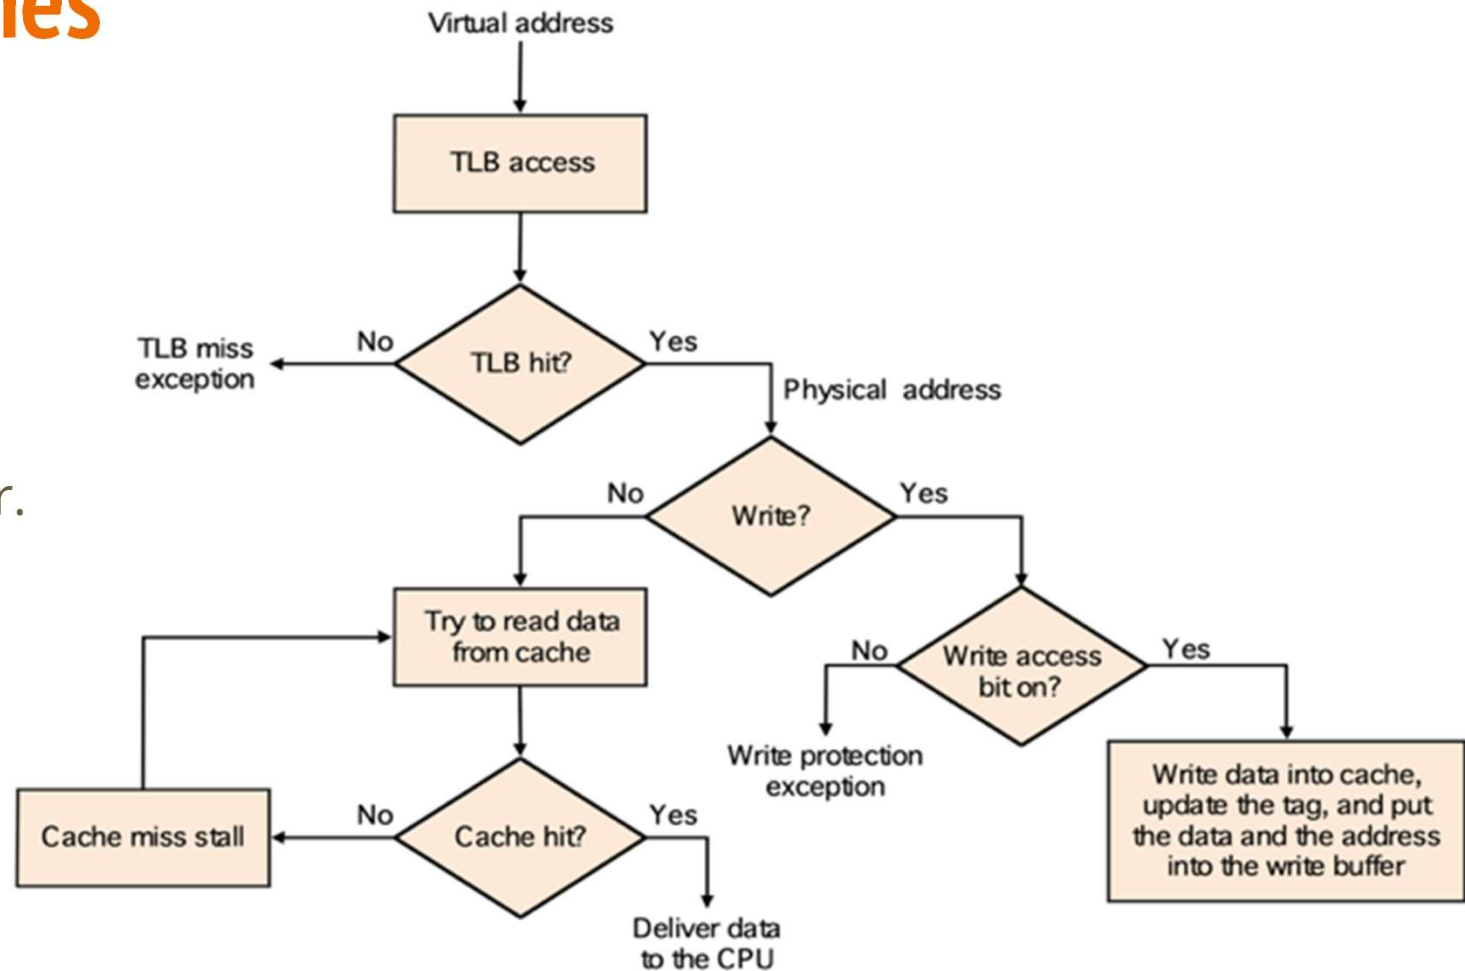


Picture from Hennessy & Patterson



TLB & Caches

- ★ Steps in accessing data (taken from Intel x86)
- ★ This is slow.
- ★ Let's make it faster.






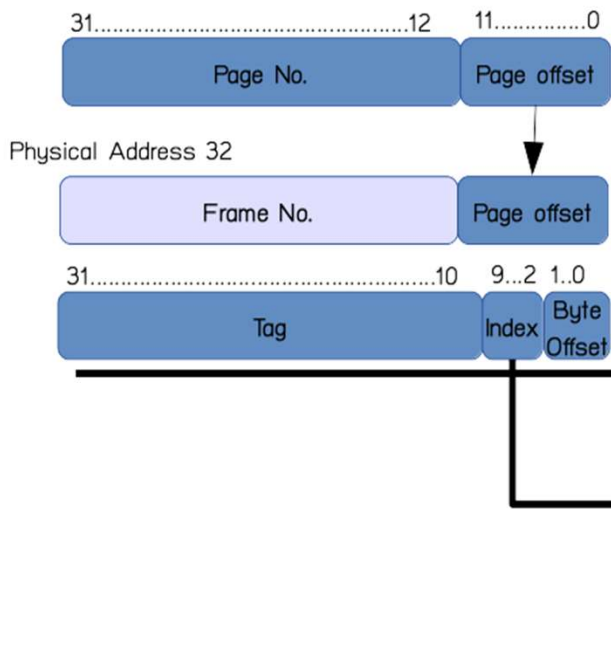
Reduce Translation Time

- ★ Machines with TLBs go one step further to reduce # cycles/cache access
- ★ They overlap the cache access with the TLB access
- ★ Works because high order bits of the VA are used to look in the TLB while low order bits are used as index into cache



Overlapped Cache & TLB Access

Virtual Address 32 bit



- ★ If we can parallel (overlap) the cache look up (index) and translation from virtual address to physical address at the same time.
- ★ Possible when **index** is **not changed** in address translation



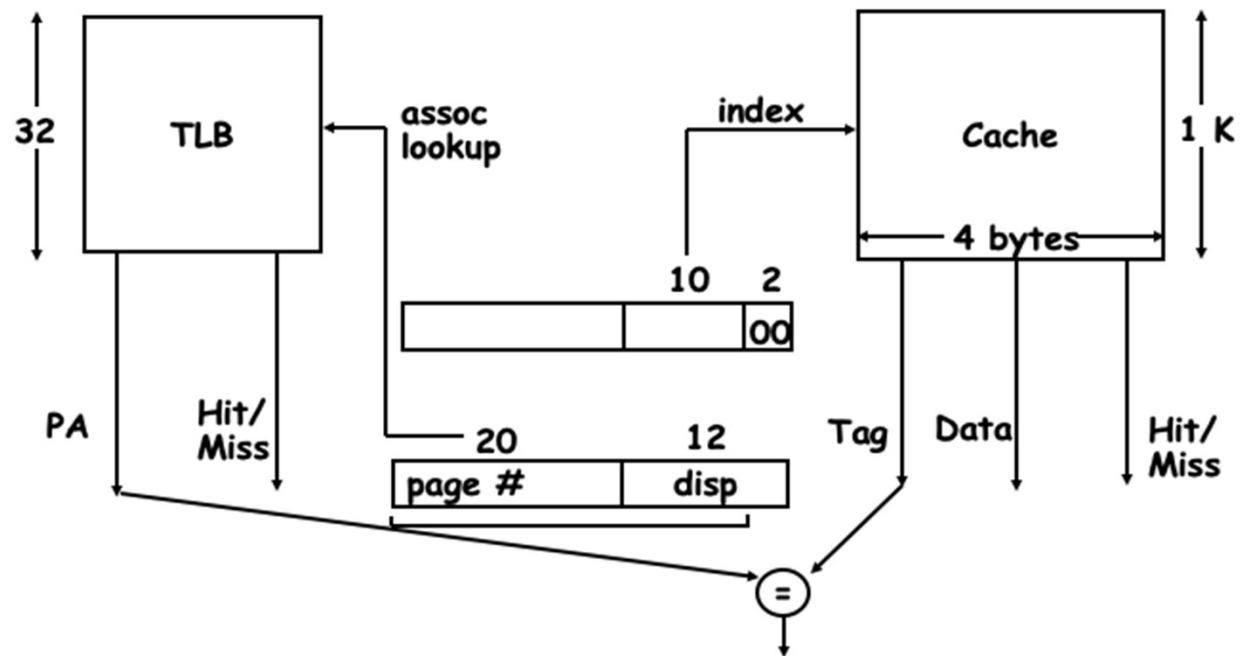
Overlapped Cache & TLB Access (ctd.)

IF cache hit AND (cache tag = PA)
then deliver data to CPU

ELSE IF [cache miss OR (cache
tag \neq PA)] and TLB hit

THEN access memory with
the PA from the TLB

ELSE do standard VA
translation






When overlapped access is not possible

- ★ Larger cache
(e.g. 8KB) in 4KB page.
- ★ 13 bits cache vs. 12 bits
page offset

Virtual Address 32 bit






When overlapped access is not possible

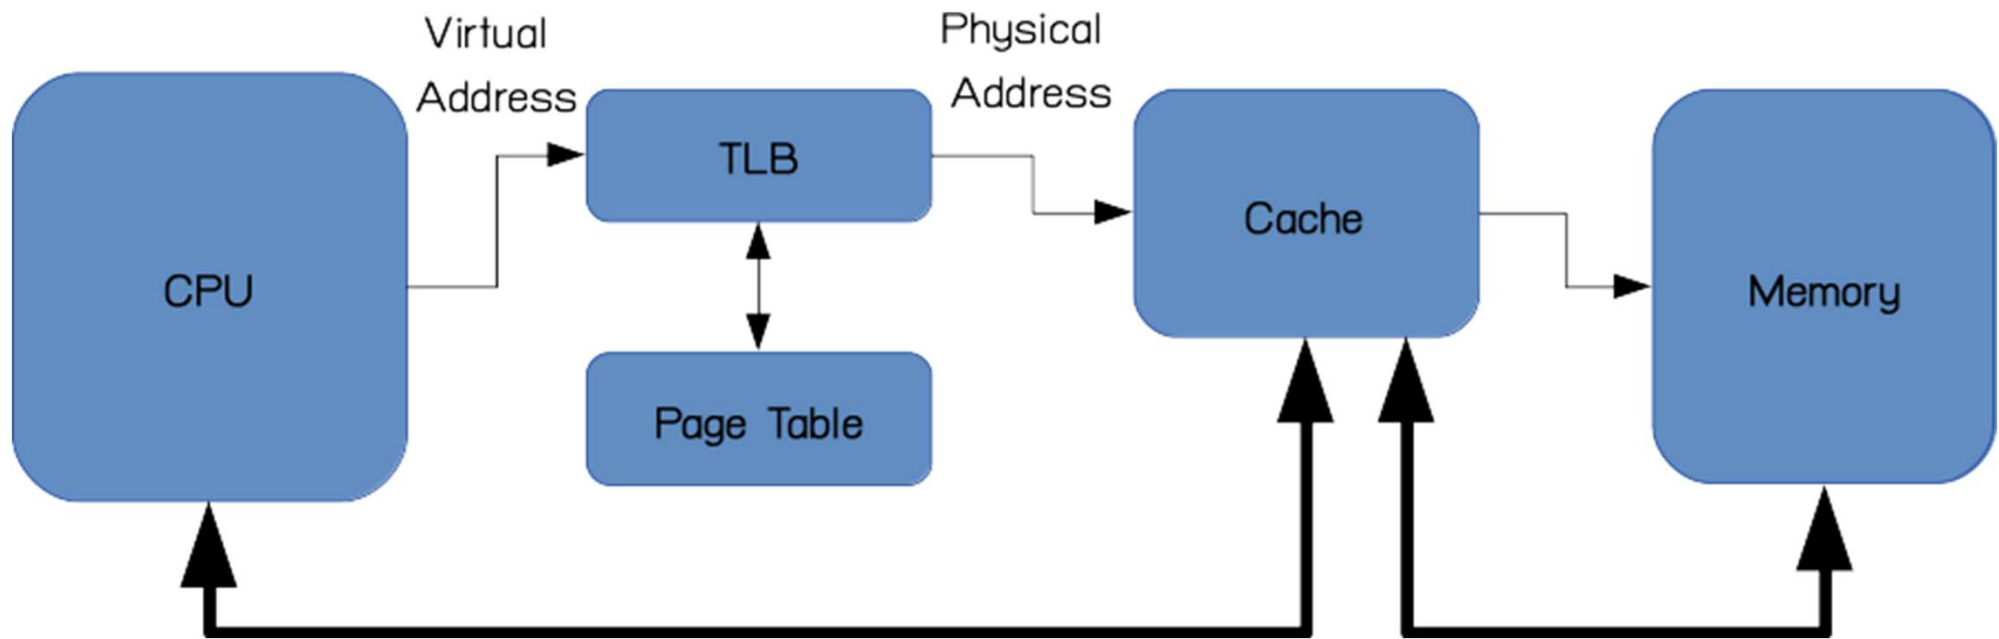
★ Solutions:

- go to 8K byte page sizes;
- go to 2 way set associative cache; or
- SW guarantee $VA[13]=PA[13]$






Virtual Memory and Cache





Can we access cache with virtual address?

- ★ Will there be any benefit (or issue) if we move cache before address translation?
- ★ Issues with shared memory (where different virtual addresses from different processes are mapped to same physical address.)





Beyond the scope of this class

- ★ There exists several technologies for performing memory relocation and memory management (such as overlay and segmentation).
- ★ There are also several methods for managing page table. (Inverted Page Table, Multilevel page table)

With virtualization, several processors now also support virtualized page table.



Summary

- ★ Caches, TLBs, Virtual Memory all understood by examining how they deal with 4 questions:
 - 1) Where can block be placed?
 - 2) How is block found?
 - 3) What block is replaced on miss?
 - 4) How are writes handled?
- ★ Page tables map virtual address to physical address
- ★ TLBs are important for fast translation
- ★ TLB misses are significant in processor performance

<https://gabrieletolomei.wordpress.com/miscellanea/operating-systems/virtual-memory-paging-and-swapping/>



Exercises





Memory Design

- ★ With 4KB page, how should an architect design 32KB cache to the advantage of overlapped access? Please provide your analysis.



Process switch

★What should happen to the following elements when the OS performs a process switch

- ★TLB
- ★Page table
- ★Cache



End of Chapter 8 (part 2)

